

R Notebook

Code ▼

Assignment 06

Baran Gulmez
2534089



Load Libraries

Hide

```
library(mlr)
library(tidyverse)
library(DataExplorer)
library(factoextra)
library(dendextend)
library(reshape2)
library(ggforce)
library(cluster)
library(corrplot)
library(ggplot2)
library(ggpubr) # this might be conflicting (select function)
library(MASS) # this might be conflicting (select function)
               # MASS library is used isoMDS()
library(conflicted)
library(factoextra) # for fviz
library(clusterCrit) # cluster validation (intCriteria function)
library(clValid) # (clValid function)
```

Read Dataset

The sample Dataset summarizes the usage behavior of about 9000 active credit card holders during the last 6 months.

CUSTID : Identification of Credit Card holder (Categorical)

BALANCE : Balance amount left in their account to make purchases

BALANCEFREQUENCY : How frequently the Balance is updated, score between 0 and 1 (1 = frequently updated, 0 = not frequently updated)

PURCHASES : Amount of purchases made from account

ONEOFFPURCHASES : Maximum purchase amount done in one-go

INSTALLMENTSPURCHASES : Amount of purchase done in installment

CASHADVANCE : Cash in advance given by the user

PURCHASESFREQUENCY : How frequently the Purchases are being made, score between 0 and 1 (1 = frequently purchased, 0 = not frequently purchased)

ONEOFFPURCHASESFREQUENCY : How frequently Purchases are happening in one-go (1 = frequently purchased, 0 = not frequently purchased)

PURCHASEINSTALLMENTSFREQUENCY : How frequently purchases in installments are being done (1 = frequently done, 0 = not frequently done)

CASHADVANCEFREQUENCY : How frequently the cash in advance being paid

CASHADVANCETRX : Number of Transactions made with "Cash in Advanced"

PURCHASESTRX : Numbe of purchase transactions made

CREDITLIMIT : Limit of Credit Card for user

PAYMENTS : Amount of Payment done by user

MINIMUM_PAYMENTS : Minimum amount of payments made by user

PRCFULLPAYMENT : Percent of full payment paid by user

TENURE : Tenure of credit card service for user

[Hide](#)

```
dat = read.csv("input/CC GENERAL.csv",  
              stringsAsFactors = F,  
              na.strings = c(" "))
```

Basic Analysis

This dataset is quite useful since because of two reasons. The first is that the dataset does not need preprocessing since all features are numeric. The second is that there are more than enough data in terms of both number of samples and number of features. ## Glimpse

[Hide](#)

```
glimpse(dat)
```

Rows: 8,950

Columns: 18

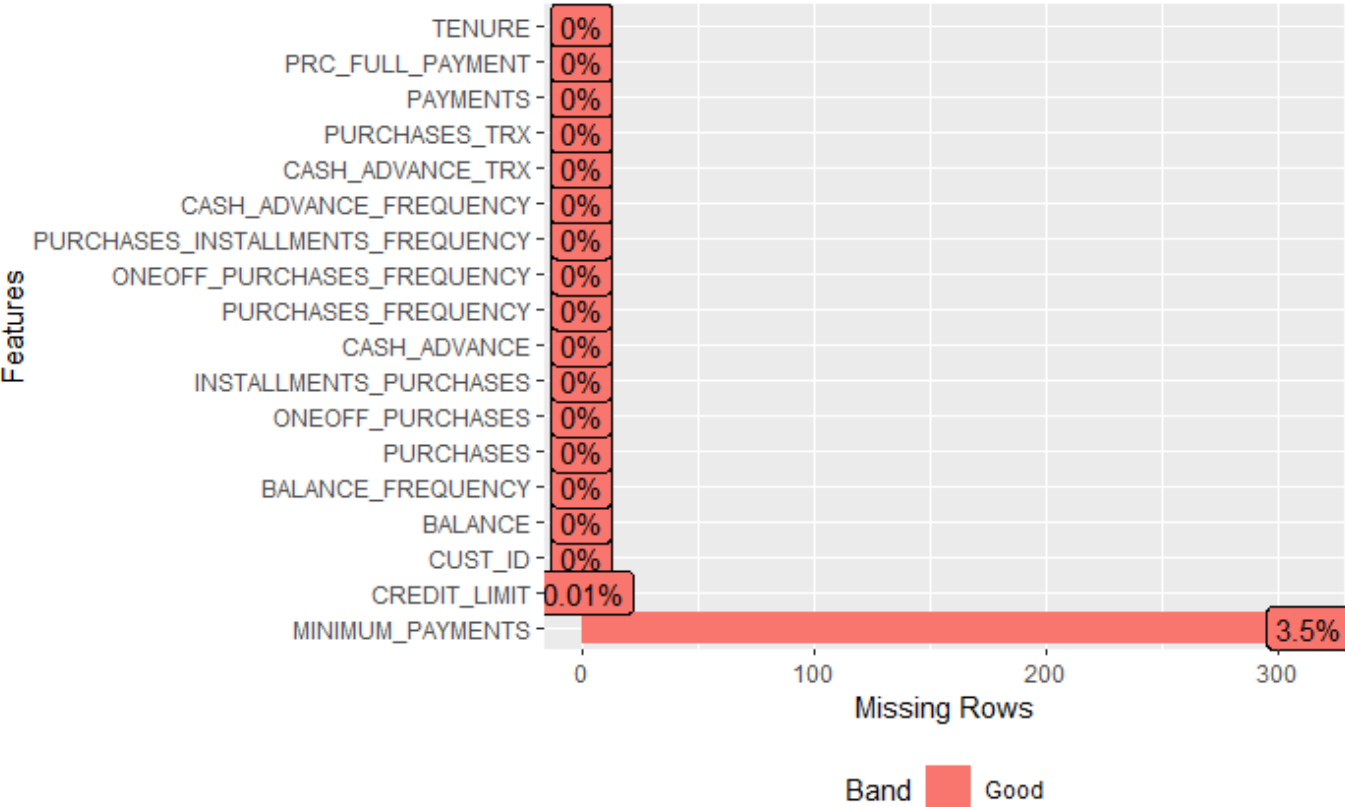
```
$ CUST_ID          <chr> "C10001", "C10002", "C10003", "C10004", "C10005", "C
10006", "C10007", "C10008", "C10009", "C10010~
$ BALANCE          <dbl> 40.90075, 3202.46742, 2495.14886, 1666.67054, 817.71
434, 1809.82875, 627.26081, 1823.65274, 1014.~
$ BALANCE_FREQUENCY <dbl> 0.818182, 0.909091, 1.000000, 0.636364, 1.000000, 1.
000000, 1.000000, 1.000000, 0.54545~
$ PURCHASES        <dbl> 95.40, 0.00, 773.17, 1499.00, 16.00, 1333.28, 7091.0
1, 436.20, 861.49, 1281.60, 920.12, 1492.18, ~
$ ONEOFF_PURCHASES <dbl> 0.00, 0.00, 773.17, 1499.00, 16.00, 0.00, 6402.63,
0.00, 661.49, 1281.60, 0.00, 1492.18, 2500.23,~
$ INSTALLMENTS_PURCHASES <dbl> 95.40, 0.00, 0.00, 0.00, 0.00, 1333.28, 688.38, 436.
20, 200.00, 0.00, 920.12, 0.00, 717.76, 1717.~
$ CASH_ADVANCE      <dbl> 0.00000, 6442.94548, 0.00000, 205.78802, 0.00000, 0.
00000, 0.00000, 0.00000, 0.00000, 0.00000, 0.~
$ PURCHASES_FREQUENCY <dbl> 0.166667, 0.000000, 1.000000, 0.083333, 0.083333, 0.
666667, 1.000000, 1.000000, 0.333333, 0.16666~
$ ONEOFF_PURCHASES_FREQUENCY <dbl> 0.000000, 0.000000, 1.000000, 0.083333, 0.083333, 0.
000000, 1.000000, 0.000000, 0.083333, 0.16666~
$ PURCHASES_INSTALLMENTS_FREQUENCY <dbl> 0.083333, 0.000000, 0.000000, 0.000000, 0.000000, 0.
583333, 1.000000, 1.000000, 0.250000, 0.00000~
$ CASH_ADVANCE_FREQUENCY <dbl> 0.000000, 0.250000, 0.000000, 0.083333, 0.000000, 0.
000000, 0.000000, 0.000000, 0.000000, 0.00000~
$ CASH_ADVANCE_TRX   <int> 0, 4, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 4, 3,
0, 0, 0, 0, 6, 0, 13, 4, 0, 5, 0, 16, 0, 10, 2~
$ PURCHASES_TRX      <int> 2, 0, 12, 1, 1, 8, 64, 12, 5, 3, 12, 6, 26, 26, 0, 1
1, 0, 8, 9, 12, 8, 92, 17, 13, 0, 12, 2, 12, ~
$ CREDIT_LIMIT        <dbl> 1000, 7000, 7500, 7500, 1200, 1800, 13500, 2300, 700
0, 11000, 1200, 2000, 3000, 7500, 3000, 8000,~
$ PAYMENTS            <dbl> 201.8021, 4103.0326, 622.0667, 0.0000, 678.3348, 140
0.0578, 6354.3143, 679.0651, 688.2786, 1164.7~
$ MINIMUM_PAYMENTS    <dbl> 139.50979, 1072.34022, 627.28479, NA, 244.79124, 240
7.24604, 198.06589, 532.03399, 311.96341, 100~
$ PRC_FULL_PAYMENT    <dbl> 0.000000, 0.222222, 0.000000, 0.000000, 0.000000, 0.
000000, 1.000000, 0.000000, 0.000000, 0.00000~
$ TENURE              <int> 12, 12, 12, 12, 12, 12, 12, 12, 12, 12, 12, 12, 12,
12, 12, 12, 8, 12, 12, 12, 12, 12, 12, 11~
```

Missing Data

An insignificant portion of the data is missing.

Hide

```
plot_missing(dat)
```



Summaryy

Hide

```
summary(dat)
```

CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES
INSTALLMENTS_PURCHASES	CASH_ADVANCE			
Length:8950	Min. : 0.0	Min. :0.0000	Min. : 0.00	Min. : 0.0
Min. : 0.0	Min. : 0.0			
Class :character	1st Qu.: 128.3	1st Qu.:0.8889	1st Qu.: 39.63	1st Qu.: 0.0
1st Qu.: 0.0	1st Qu.: 0.0			
Mode :character	Median : 873.4	Median :1.0000	Median : 361.28	Median : 38.0
Median : 89.0	Median : 0.0			
	Mean : 1564.5	Mean :0.8773	Mean : 1003.20	Mean : 592.4
Mean : 411.1	Mean : 978.9			
	3rd Qu.: 2054.1	3rd Qu.:1.0000	3rd Qu.: 1110.13	3rd Qu.: 577.4
3rd Qu.: 468.6	3rd Qu.: 1113.8			
	Max. :19043.1	Max. :1.0000	Max. :49039.57	Max. :40761.2
Max. :22500.0	Max. :47137.2			

PURCHASES_FREQUENCY	ONEOFF_PURCHASES_FREQUENCY	PURCHASES_INSTALLMENTS_FREQUENCY	CASH_ADVANCE_FREQUENCY
CASH_ADVANCE_TRX	PURCHASES_TRX		
Min. :0.00000	Min. :0.00000	Min. :0.0000	Min. :0.00
00	Min. : 0.000	Min. : 0.00	
1st Qu.:0.08333	1st Qu.:0.00000	1st Qu.:0.0000	1st Qu.:0.00
00	1st Qu.: 0.000	1st Qu.: 1.00	
Median :0.50000	Median :0.08333	Median :0.1667	Median :0.00
00	Median : 0.000	Median : 7.00	
Mean :0.49035	Mean :0.20246	Mean :0.3644	Mean :0.13
51	Mean : 3.249	Mean : 14.71	
3rd Qu.:0.91667	3rd Qu.:0.30000	3rd Qu.:0.7500	3rd Qu.:0.22
22	3rd Qu.: 4.000	3rd Qu.: 17.00	
Max. :1.00000	Max. :1.00000	Max. :1.0000	Max. :1.50
00	Max. :123.000	Max. :358.00	


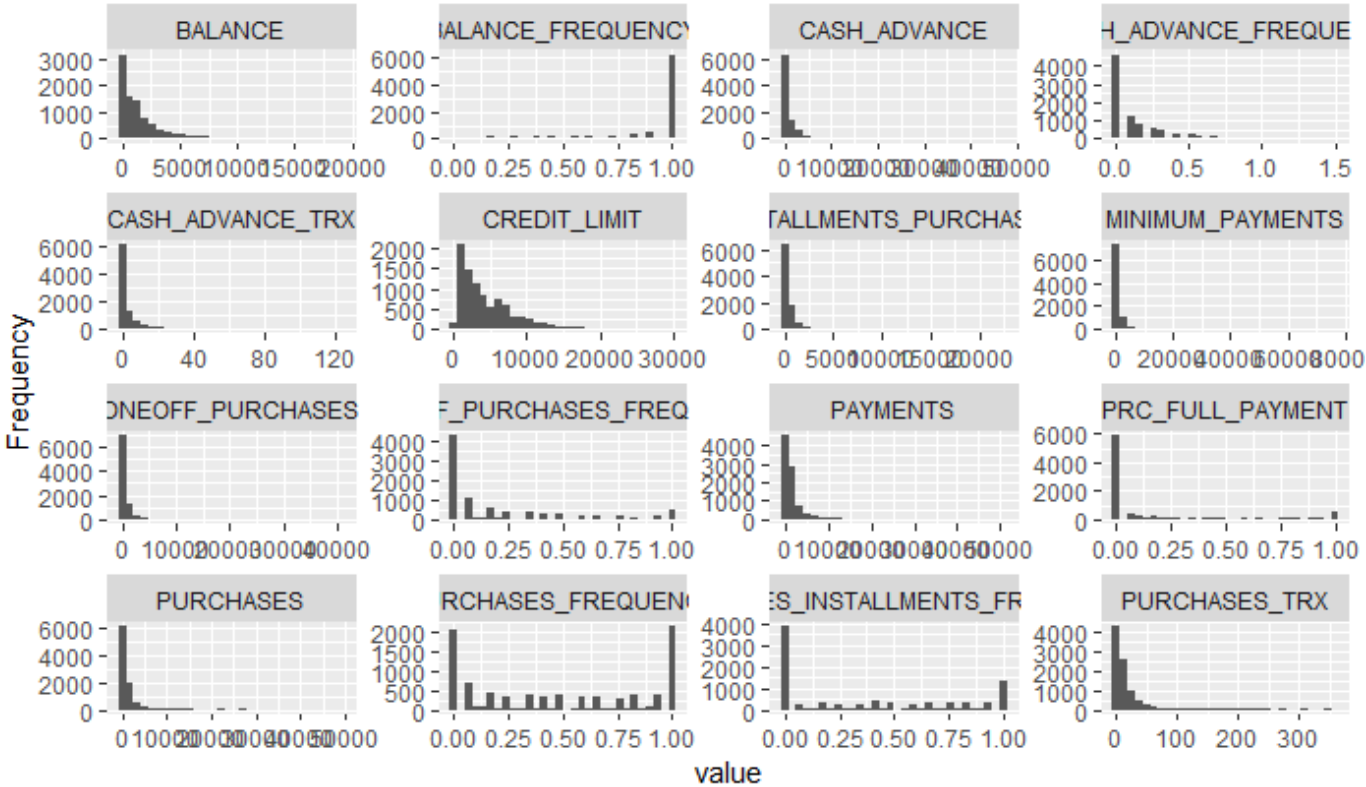
CREDIT_LIMIT	PAYMENTS	MINIMUM_PAYMENTS	PRC_FULL_PAYMENT	TENURE
Min. : 50	Min. : 0.0	Min. : 0.02	Min. :0.0000	Min. : 6.00
1st Qu.: 1600	1st Qu.: 383.3	1st Qu.: 169.12	1st Qu.:0.0000	1st Qu.:12.00
Median : 3000	Median : 856.9	Median : 312.34	Median :0.0000	Median :12.00
Mean : 4494	Mean : 1733.1	Mean : 864.21	Mean :0.1537	Mean :11.52
3rd Qu.: 6500	3rd Qu.: 1901.1	3rd Qu.: 825.49	3rd Qu.:0.1429	3rd Qu.:12.00
Max. :30000	Max. :50721.5	Max. :76406.21	Max. :1.0000	Max. :12.00
NA's :1		NA's :313		

Histograms

Here it is seen that almost all features are skewed.

Hide

```
plot_histogram(dat)
```



Data Reorganization

Missing data deleted and some useless CUST_ID deleted.

Hide

```

dat_reorg = dat %>%
  dplyr::select(-CUST_ID) %>% # calling select from dplyr to prevent conflict
  drop_na()
#
# shuffle data
if(FALSE){
  set.seed(357) # fix seed
  dat_reorg_backup <- dat_reorg[sample(nrow(dat_reorg)), ] # generate random index using sample
  and
  dat_reorg = dat_reorg_backup[1:100,1:ncol(dat_reorg_backup)] # choose first 100
}

```


Correlation

[Hide](#)

```

corrplot(cor(dat_reorg), diag = FALSE, type = "lower", tl.srt = 45, tl.col = "black", method
= 'color', tl.cex = 0.4)

```



PCA

prcomp() expects the samples to be rows to be columns

[Hide](#)

```

dat_reorg_scaled = scale(dat_reorg)
pca <- prcomp(dat_reorg_scaled)


```

PCA Analysis

Scree plot shows how much the principal components are responsible the PCA component is responsible of the variation of the data.

Hide

```
pca.var <- pca$sdev^2
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)
barplot(pca.var.per, main="Scree Plot", xlab="Principal Component", ylab="Percent Variation")
```



format the data

Hide

```
pca.data <- data.frame(Sample=rownames(dat_reorg), X=pca$x[,1], Y=pca$x[,2])
pca.data
```

Sample <chr>	X <dbl>	Y <dbl>
1	-1.696297065	-1.122518989
2	-1.215610445	2.435496753
3	0.935799104	-0.385179263
4	-1.614544792	-0.724544205
5	0.223687663	-0.783564452
6	6.265235048	-0.609413897
7	0.261651730	-1.295560288
8	-0.465311641	-0.477670250

	BALANCE	BALANCE_FREQUENCY	PURC
HASES	ONEOFF_PURCHASES		
	0.09198590	0.10981218	0.412
15123	0.34677536		
	INSTALLMENTS_PURCHASES	CASH_ADVANCE	PURCHASES_FREQ
UENCY	ONEOFF_PURCHASES_FREQUENCY		
	0.33705564	-0.03058765	0.323
66488	0.29476135		
PURCHASES_INSTALLMENTS_FREQUENCY		CASH_ADVANCE_FREQUENCY	CASH_ADVANC
E_TRX	PURCHASES_TRX		
	0.27722626	-0.09914541	-0.056
96036	0.39106653		
	CREDIT_LIMIT	PAYMENTS	MINIMUM_PAY
MENTS	PRC_FULL_PAYMENT		
	0.21005184	0.26372547	0.059
32632	0.13056503		
	TENURE		
	0.07791867		

Multi-Dimensional Scaling (MDS)

Euclidean and manhattan distances are most widely used distance metrics. Therefore they are chosen as distance metrics. These two distances also usually work the best to my experience.

Hide

```
distEuc.matrix <- stats::dist(dat_reorg_scaled, method="euclidean")
distMnh.matrix <- stats::dist(dat_reorg_scaled, method="manhattan")
```

Classical Multi-Dimensional Scaling

Euclidean Distance

Hide


```
mdsCmdEuc.stuff <- cmdscale(distEuc.matrix, eig=TRUE, x.ret=TRUE)
mdsCmdEuc.var.per <- round(mdsCmdEuc.stuff$eig/sum(mdsCmdEuc.stuff$eig)*100, 1)
# graph
mdsCmdEuc.values <- mdsCmdEuc.stuff$points
mdsCmdEuc.data <- base::data.frame(Sample=rownames(dat_reorg), X=mdsCmdEuc.values[,1], Y=mdsC
mdEuc.values[,2])
mdsCmdEuc.data
```

Sample	X	Y
<chr>	<dbl>	<dbl>
1	1.696297065	1.122518989
2	1.215610445	-2.435496753
3	-0.935799104	0.385179263
4	1.614544792	0.724544205
5	-0.223687663	0.783564452

Sample	X	Y
<chr>	<dbl>	<dbl>
6	-6.265235048	0.609413897
7	-0.261651730	1.295560288
8	0.465311641	0.477670250
9	0.599646437	0.408567662
10	-0.522740791	1.312087280
1-10 of 8,636 rows		
	Previous	1 2 3 4 5 6 ... 100 Next

Hide


```
ggplot(data=mdsCmdEuc.data, aes(x=X, y=Y, label=Sample)) +
  geom_text() +
  theme_bw() +
  xlab(paste("MDS1 - ", mdsCmdEuc.var.per[1], "%", sep="")) +
  ylab(paste("MDS2 - ", mdsCmdEuc.var.per[2], "%", sep=""))
```



Manhattan Distance

Hide

```
mdsCmdMnh.stuff <- cmdscale(distMnh.matrix, eig=TRUE, x.ret=TRUE)
mdsCmdMnh.var.per <- round(mdsCmdMnh.stuff$eig/sum(mdsCmdMnh.stuff$eig)*100, 1)
# graph
mdsCmdMnh.values <- mdsCmdMnh.stuff$points
mdsCmdMnh.data <- data.frame(Sample=rownames(dat_reorg),
  X=mdsCmdMnh.values[,1],
  Y=mdsCmdMnh.values[,2])
ggplot(data=mdsCmdMnh.data, aes(x=X, y=Y, label=Sample)) +
  geom_text() +
  theme_bw() +
  xlab(paste("MDS1 - ", mdsCmdMnh.var.per[1], "%", sep="")) +
  ylab(paste("MDS2 - ", mdsCmdMnh.var.per[2], "%", sep=""))
```



Metric Multi-Dimensional Scaling

Euclidean Distance

Hide

```
mdsSamEuc.stuff <- sammon(distEuc.matrix)
```

```
Initial stress      : 0.18811
stress after 0 iters: 0.18811
```


Hide

```
mdsSamEuc.var.per <- round(mdsSamEuc.stuff$eig/sum(mdsSamEuc.stuff$eig)*100, 1)
# graph
mdsSamEuc.values <- mdsSamEuc.stuff$points
mdsSamEuc.data <- base::data.frame(Sample=rownames(dat_reorg), X=mdsSamEuc.values[,1], Y=mdsSamEuc.values[,2])
mdsSamEuc.data
```

Sample <chr>	X <dbl>	Y <dbl>
1	1.696297065	1.122518989
2	1.215610445	-2.435496753
3	-0.935799104	0.385179263
4	1.614544792	0.724544205
5	-0.223687663	0.783564452
6	-6.265235048	0.609413897
7	-0.261651730	1.295560288
8	0.465311641	0.477670250
9	0.599646437	0.408567662
10	-0.522740791	1.312087280
1-10 of 8,636 rows		Previous 1 2 3 4 5 6 ... 100 Next

Hide

```
ggplot(data=mdsSamEuc.data, aes(x=X, y=Y, label=Sample)) +
  geom_text() +
  theme_bw() +
  xlab(paste("MDS1 - ", mdsSamEuc.var.per[1], "%", sep="")) +
  ylab(paste("MDS2 - ", mdsSamEuc.var.per[2], "%", sep=""))
```



Manhattan Distance


[Hide](#)

```
mdsSamMnh.stuff <- sammon(distMnh.matrix)
```

```
Initial stress      : 0.08243
stress after  0 iters: 0.08243
```

[Hide](#)

```
mdsSamMnh.var.per <- round(mdsSamMnh.stuff$eig/sum(mdsSamMnh.stuff$eig)*100, 1)
# graph
mdsSamMnh.values <- mdsSamMnh.stuff$points
mdsSamMnh.data <- data.frame(Sample=rownames(dat_reorg),
  X=mdsSamMnh.values[,1],
  Y=mdsSamMnh.values[,2])
ggplot(data=mdsSamMnh.data, aes(x=X, y=Y, label=Sample)) +
  geom_text() +
  theme_bw() +
  xlab(paste("MDS1 - ", mdsSamMnh.var.per[1], "%", sep="")) +
  ylab(paste("MDS2 - ", mdsSamMnh.var.per[2], "%", sep=""))
```



Non-Metric Multi-Dimensional Scaling

Euclidean Distance

Hide

```
mdsIsoEuc.stuff <- isoMDS(distEuc.matrix)
```

```
initial value 28.346012
iter 5 value 17.743101
iter 10 value 14.923771
final value 14.886499
converged
```

Hide

```
# graph
mdsIsoEuc.values <- mdsIsoEuc.stuff$points
mdsIsoEuc.data <- data.frame(Sample=rownames(dat_reorg),
  X=mdsIsoEuc.values[,1],
  Y=mdsIsoEuc.values[,2])
mdsIsoEuc.data
```

Sample	X	Y
<chr>	<dbl>	<dbl>
1	1.280355162	0.706059168
2	1.053834590	-1.813733180
3	-1.370248828	-0.140027266


```
initial value 22.024305
final value 22.020486
converged
```


Hide

```
# graph
mdsIsoMnh.values <- mdsIsoMnh.stuff$points
mdsIsoMnh.data <- data.frame(Sample=rownames(dat_reorg),
  X=mdsIsoMnh.values[,1],
  Y=mdsIsoMnh.values[,2])
mdsIsoMnh.data
```

Sample <chr>	X <dbl>	Y <dbl>
1	-4.12217468	-4.636481458
2	-7.06373977	5.807594099
3	3.85313180	1.463960055
4	-3.71610128	-2.772121054
5	2.29220920	-0.495867094
6	22.81053019	-1.216820052
7	2.96938601	-1.460609252
8	-0.12917346	-0.424079563
9	-1.86588575	-2.271565354
10	3.80709155	-1.360532928
1-10 of 8,636 rows		Previous 1 2 3 4 5 6 ... 100 Next

Hide

```
ggplot(data=mdsIsoMnh.data, aes(x=X, y=Y, label=Sample)) +
  geom_text() +
  theme_bw() +
  xlab(paste("MDS1 - ")) +
  ylab(paste("MDS2 - "))
```



MDS Results

There is not much difference between euclidean and manhattan distances. Both Classical MDS and Non-Metric MDS support this. When Classical MDS and Non-Metric MDS are compared, classical MDS separates the samples into a wider spectrum which is better. This will help get more distinct groups during clustering.


Clustering

The reason for me to separate the data into 3 clusters is totally intuitively. I just assumed that people would be from low, mid and high income. But inspecting hierarchical clustering, there could be 4-6 clusters ideally.

Hierarchical


Hide

```
clustHrc = hclust(distEuc.matrix , method = "ward.D")
plot(clustHrc, labels = FALSE, sub = "", xlab = "", ylab = "Euclidian Dist")
rect.hclust(clustHrc, k = 3)
```



Hide


```
clustHrcClusters = cutree(clustHrc, k = 3)
fviz_pca_ind(pca, habillage = clustHrcClusters)
```




K-Means

Hide

```
# draw optimal number of cluster
fviz_nbclust(dat_reorg_scaled, kmeans, method = "wss", k.max = 10)
```


[Hide](#)

```
clustKm = kmeans(dat_reorg_scaled, centers = 3)
fviz_pca_ind(pca, habillage = clustKm$cluster)
```



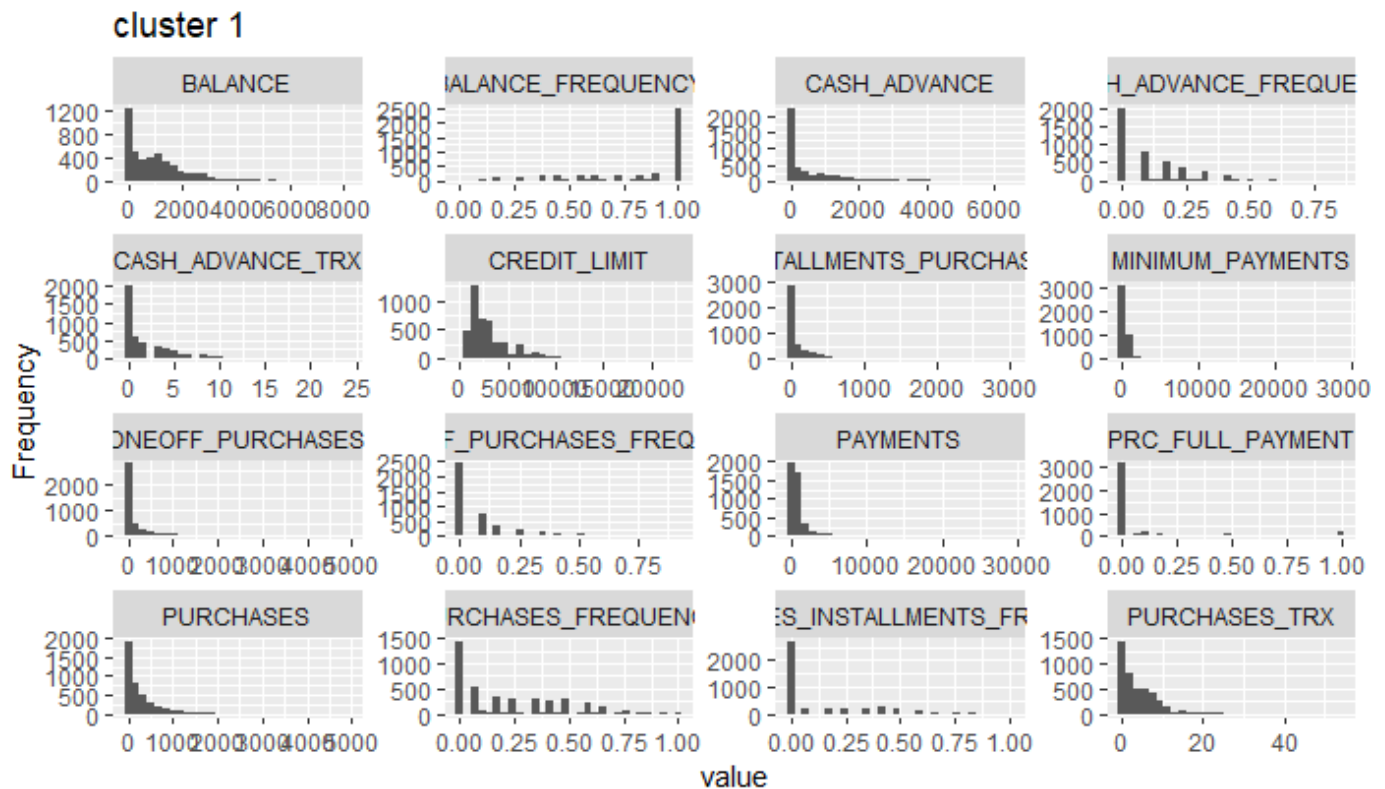
Analyze Groups

[Hide](#)


```

clustKm1.index = which(clustKm$cluster==1)
clustKm2.index = which(clustKm$cluster==2)
clustKm3.index = which(clustKm$cluster==3)
clustKm1.dat = dat_reorg[clustKm1.index,1:ncol(dat_reorg)]
clustKm2.dat = dat_reorg[clustKm2.index,1:ncol(dat_reorg)]
clustKm3.dat = dat_reorg[clustKm3.index,1:ncol(dat_reorg)]
plot_histogram(clustKm1.dat, title="cluster 1")

```



Page 1



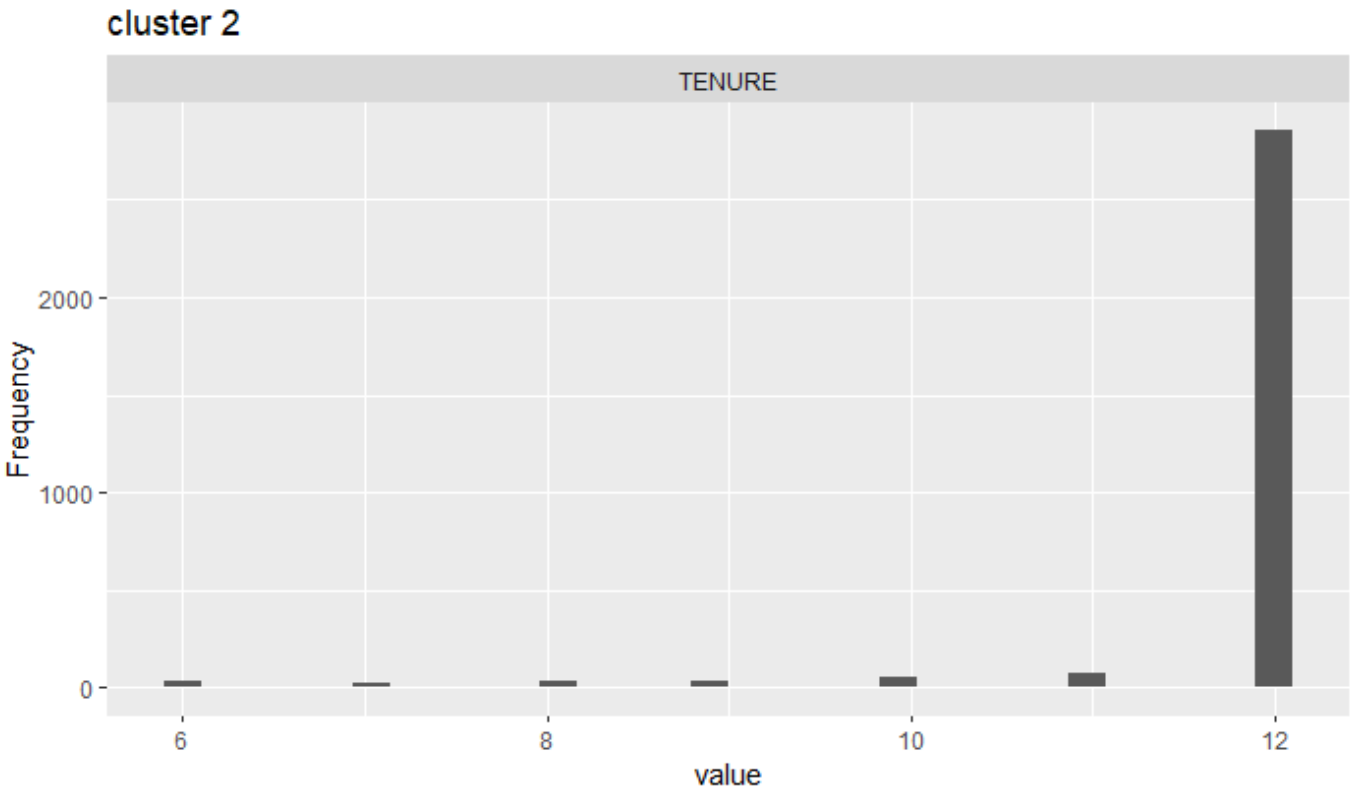
Page 2

Hide

```
plot_histogram(clustKm2.dat, title="cluster 2")
```



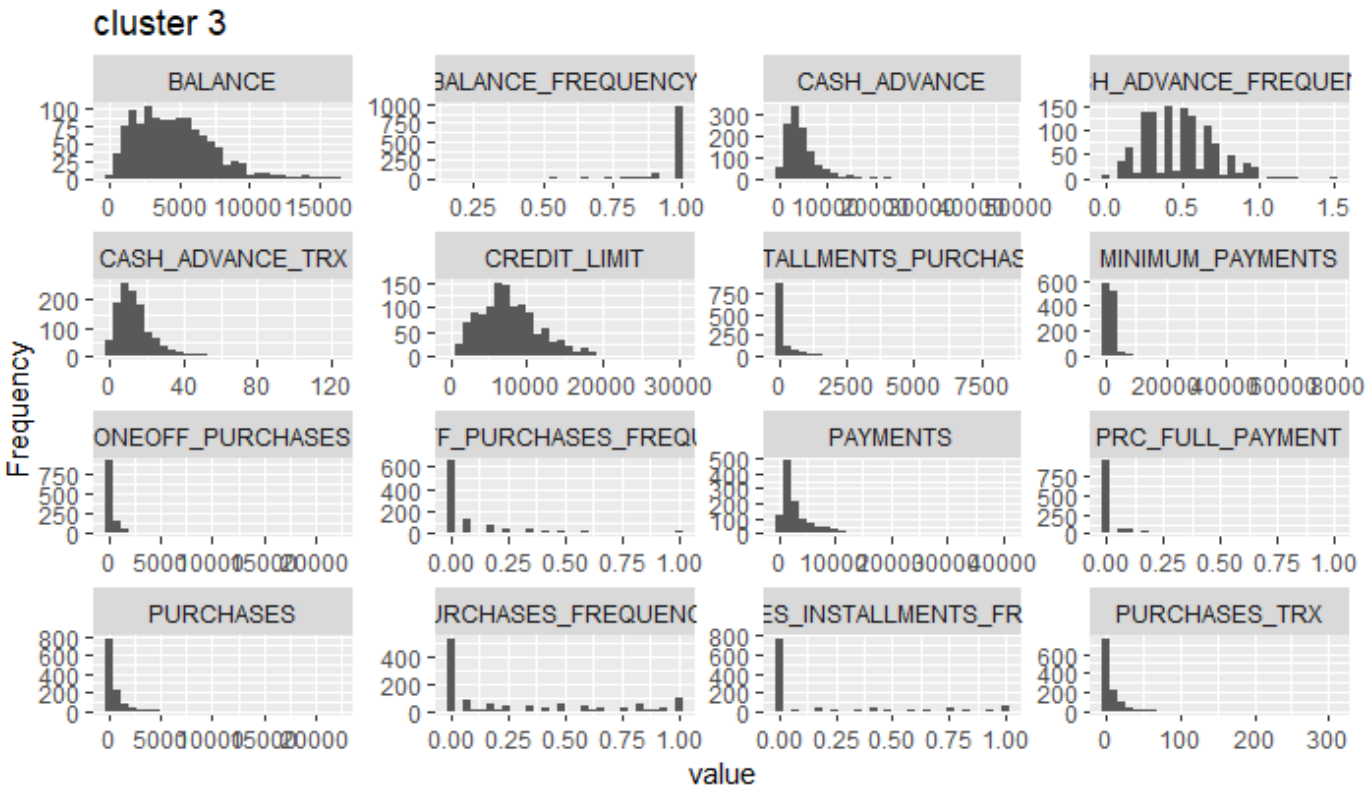
Page 1




Page 2

Hide

```
plot_histogram(clustKm3.dat, title="cluster 3")
```



Page 1



Page 2

After clustering the samples I inspected each cluster separately. According to my intuitive assumption about having 3 income groups, the groups are: "rich", "middle class" and "poor". When the group 1 is inspected they have higher credit limit, balance and have higher purchase frequency. This and other features clearly reveals that this group is the rich one. Accordingly group 2 is the "middle class" and group 3 is the "poor" group. # Validation Dunn Index(DI), Davies-Bouldin(DBI) Index and Silhouette Coefficient are inspected.

Clustering gets better as Dunn Index increases. DI evaluates the clusters using the farthest points and in this dataset there are very far points which I think outliers. DI could be more meaningful if the dataset would not contain outliers. DI decreases as the number of clusters increase thus it indicates number of clusters being lower is better.

Clustering gets better as Davies-Bouldin Index decreases. DBI is a metric of separation of clusters. When scores are inspected, hierarchical clustering gives much better DBI(Connectivity) scores than k-means clustering. So, hierarchical clustering is much better at separating clusters. This can also be visually seen when the colored clusters are inspected.

Clustering gets better as Silhouette Coefficient increases. Since its range is between 0 and 1 the optimal value is 1. SC is expected to be higher than 0.5 and again hierarchical clustering is much better.

Even though we do these validation measures after clustering, in practice they can be applied before clustering in order to understand which configuration should be used.

Hide

```
# https://rdrr.io/cran/clusterCrit/man/intCriteria.html
intCriteria(dat_reorg_scaled, clustKm$cluster, c("Dunn", "dav", "silhouette") )
```



```
$dunn
[1] 0.002012817

$davies_bouldin
[1] 1.663439

$silhouette
[1] 0.1524162
```

Internal Validation

Hide

```
# https://www.rdocumentation.org/packages/clValid/versions/0.7/topics/clValid
# https://cran.r-project.org/web/packages/clValid/vignettes/clValid.pdf
# https://rdr.io/cran/clValid/man/clValid-class.html
valid.intern <- clValid(dat_reorg , 2:6, clMethods=c("hierarchical","kmeans"), validation="internal")
```

The number of items to be clustered is larger than 'maxitems'
The memory and time required may be excessive, do you wish to continue?
(y to continue, any other character to exit)

Hide

```
y
summary(valid.intern)
```

Clustering Methods:						
hierarchical kmeans						
Cluster sizes:						
2 3 4 5 6						
Validation Measures:						
		2	3	4	5	6
hierarchical	Connectivity	3.2956	6.2246	14.4980	19.7726	31.1563
	Dunn	0.1943	0.1943	0.1020	0.1020	0.1020
	Silhouette	0.9081	0.8842	0.8526	0.8433	0.7935
kmeans	Connectivity	555.7802	546.5504	650.4917	855.3817	965.4060
	Dunn	0.0019	0.0025	0.0021	0.0048	0.0023
	Silhouette	0.5113	0.5095	0.4628	0.3930	0.3743
Optimal Scores:						

	Score	Method	Clusters
	<dbl>	<chr>	<chr>
Connectivity	3.2956	hierarchical	2

	Score	Method	Clusters
	<dbl>	<chr>	<chr>
Dunn	0.1943	hierarchical	2
Silhouette	0.9081	hierarchical	2

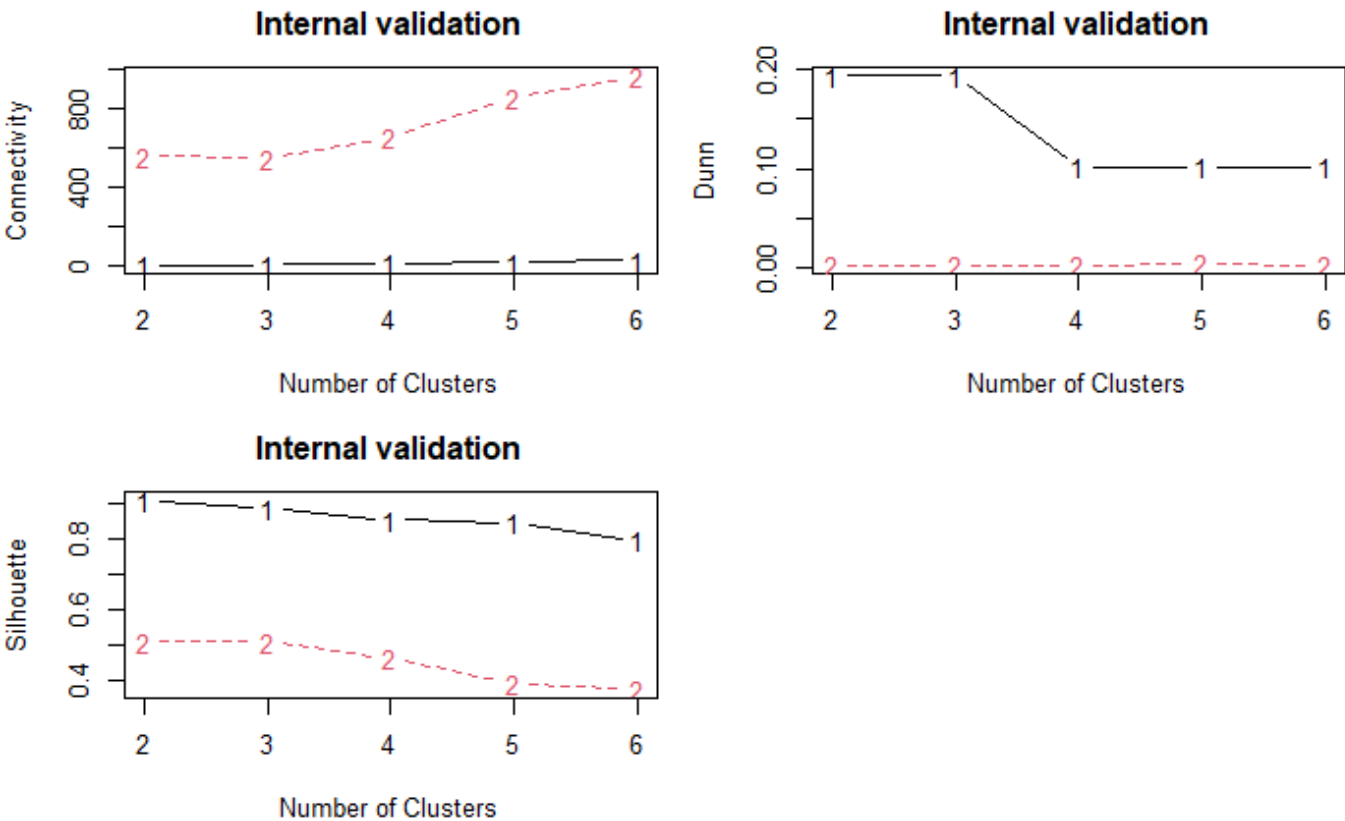
3 rows

Hide

```
op <- par(no.readonly=TRUE)
par(mfrow=c(2,2),mar=c(4,4,3,1))
plot(valid.intern, legend=FALSE)
plot(nClusters(valid.intern), measures(valid.intern,"Dunn")[,1],type="n",axes=F,xlab="",ylab="" )
```

Hide

```
legend("center", clusterMethods(valid.intern), col=1:9, lty=1:9, pch=paste(1:9))
par(op)
```



Above results show that optimal scores are achieved by hierarchical clustering with k=2.

Stability Validation

Stability validation validates reproducibility of clustering solution another sample.The included measures are the average proportion of non-overlap (APN), the average distance (AD), the average distance between means (ADM), and the figure of merit (FOM) (Datta and Datta, 2003; Yeung et al., 2001). . In all cases the average is taken over all the deleted columns, and all measures should be minimized.

Hide

```
valid.stab <- clValid(dat_reorg , 2:6, clMethods=c("hierarchical","kmeans"), validation="stability")
```

The number of items to be clustered is larger than 'maxitems'
The memory and time required may be excessive, do you wish to continue?
(y to continue, any other character to exit)

Hide

```
y  
summary(valid.stab)
```

Clustering Methods:
hierarchical kmeans

Cluster sizes:
2 3 4 5 6

Validation Measures:

		2	3	4	5	6
hierarchical	APN	0.0001	0.0009	0.0005	0.0013	0.0020
	AD	7005.3751	6992.6928	6832.9555	6825.3289	6801.3493
	ADM	7.4382	41.5973	31.1390	55.5458	77.1517
	FOM	1056.3065	1043.0755	990.3924	990.0132	988.3864
kmeans	APN	0.0221	0.0298	0.0344	0.0538	0.0670
	AD	5738.2163	5618.7797	5282.4759	4990.7224	4839.7741
	ADM	257.3348	362.1208	361.9208	561.9007	663.1631
	FOM	985.5752	991.3582	941.7552	936.3058	912.5554

Optimal Scores:

	Score	Method	Clusters
	<dbl>	<chr>	<chr>
APN	0.0001	hierarchical	2
AD	4839.7741	kmeans	6
ADM	7.4382	hierarchical	2
FOM	912.5554	kmeans	6


4 rows

Hide

```
par(mfrow=c(2,2),mar=c(4,4,3,1))  
plot(valid.stab, measure=c("APN","AD","ADM"),legend=FALSE)  
plot(nClusters(valid.stab),measures(valid.stab,"APN")[,1],type="n",axes=F,xlab="",ylab="")
```

Hide

```
legend("center", clusterMethods(valid.stab), col=1:9, lty=1:9, pch=paste(1:9))
par(op)
```



Here we see that hierarchical-2 and kmeans-6 performs the best in terms of stability. But in internal validation hierarchical 2 was superior in all 3 metrics. Thus we can conclude that hierarchical-2 is the best.

[Hide](#)

```
valid.stabRankWeights <- getRanksWeights(valid.stab)
print(valid.stabRankWeights$ranks[,1:3], quote=FALSE)
```

	1	2	3
APN	hierarchical-2	hierarchical-4	hierarchical-3
AD	kmeans-6	kmeans-5	kmeans-4
ADM	hierarchical-2	hierarchical-4	hierarchical-3
FOM	kmeans-6	kmeans-5	kmeans-4

Conclusion

In each section comments that are related to that section are made. So in the conclusion general comments are presented. The dataset is hard to work with since the data is not best suitable for clustering. Clusters almost overlap and first two components of PCA only cover the %47(27+20) of the data. For example, if first two components would add up to %80 percent of the data we could see much distinct and non overlapping clusters. Since my dataset also does not have labels, there is no way to verify if the clustering is correct. Thus, this dataset is rather open to comment.